

Ask The Storytell AI - Technical Documentation

Project Overview

A **Retrieval-Augmented Generation (RAG)** chatbot that answers questions about classic literature (*Alice in Wonderland*, *Gulliver's Travels*, *Arabian Nights*) with witty responses, AI-generated illustrations, and audio narration. Built for an AI internship project submission.

Technology Stack

Backend

Framework: FastAPI (Python 3.11+)

Why: Async support, type safety, auto-docs, best Python AI/ML integration

Frontend

Framework: React 18 + Vite

Why: Fast HMR, modern DX, component architecture, rich ecosystem

AI Models

LLM: Google Gemini 2.0 Flash Experimental

Free tier, 2-3s responses, excellent at witty tone

Embeddings: Sentence Transformers (all-MiniLM-L6-v2)

Open source, CPU-friendly, 384-dim vectors, zero API costs

Images: Pollinations.ai (Flux model)

Free, no auth, 1-4s generation, whimsical art style

Audio: ElevenLabs TTS + OpenAI Whisper STT

Natural narration, local transcription, multilingual

Data & Search

Vector Search: Custom NumPy cosine similarity

In-memory, <50ms retrieval, disk caching

PDF Processing: PyPDF2 → custom chunking (1000/200 overlap)

Application Flow

User Query → Frontend (React)
↓
POST /api/chat → Backend (FastAPI)
↓
Embed query → SentenceTransformers
↓
Cosine similarity search → Top-5 chunks (threshold 0.25)
↓
Build prompt with context + witty persona
↓
Generate text → Gemini (temp 0.9, 800 tokens)
↓
Parallel async:

- *Generate image → Pollinations (scene from answer)*
- *Generate audio → ElevenLabs (narration)*

↓
Normalize URLs to absolute ([http://localhost:9000/static/...](http://localhost:9000/static/))
↓
Update session conversation history (last 6 messages)
↓
Return JSON → Frontend renders message + image + audio

Key Design Decisions

1. Chunking: 1000 chars, 200 overlap

Why: Balance speed, coherence, coverage

Result: 2163 chunks from 3 PDFs in ~5 seconds

2. Retrieval: Top-5, threshold 0.25

Why: Optimal context without noise; reject out-of-domain queries

Fallback: Funny "I don't know" response with image

3. Tone Control: Witty persona + temp 0.9

Prompt: "Sarcastic storyteller, spill tea ☕, 2-4 sentences max"

Result: Consistent humor across all responses

4. Multilingual: Native LLM translation

Why: Preserves wit/sarcasm better than Google Translate

How : "CRITICAL: Respond ENTIRELY in {language}" in prompt

Supports: 10 languages

5. Image Prompts: Scene from answer

Why: More relevant than keyword matching

How: Extract first 2 sentences + add "whimsical storybook art" style

6. Performance: Async + caching

Async: Image/audio generated in parallel (6s vs 9s)

Cache: Embeddings saved to disk (8s cold start vs 35s)

7. Session Memory: Last 6 messages

Why: Enable follow-up questions with context

Limit: Prevents token overflow, keeps recent context

8. Absolute URLs: Prevent broken media

Why: Frontend/backend on different ports → relative URLs fail

How: Backend returns `http://localhost:9000/static/images/...`

9. Model Switching: Centralized config

File: `config.py`

Change: LLM_PROVIDER="gemini" → "openai", zero code changes

Performance Metrics

Metric	Value	Optimization
Cold start	8-12s	Cached embeddings
Query response	5-8s	Async parallel generation
Text generation	2-3s	Gemini Flash
Image generation	1-4s	Pollinations
Retrieval	<50ms	NumPy cosine similarity
Memory	<500MB	CPU-only, in-memory
Knowledge base	2163 chunks	3 PDFs

Evaluation Criteria Met

Core Requirements

-  *Knowledge Training: PDF → PyPDF2 → chunks → embeddings → disk cache*
-  *Knowledge Retrieval: Cosine similarity, Top-5, threshold 0.25*
-  *Tone Control: Witty persona prompt, temp 0.9, emoji restraint*
-  *Image Creation: Pollinations API, scene prompts, local caching*
-  *Model Switching: Centralized `config.py`, zero code changes*
-  *Accuracy: RAG grounding, source citations, out-of-domain detection*
-  *Open Source: 90% free/open source (SentenceTransformers, FastAPI, React, Whisper)*
-  *System Design: Modular separation, async/await, error handling, session management*

Bonus Features

-  *Audio Input: Whisper STT, mic button, auto-transcribe & send*
-  *Audio Output: ElevenLabs TTS, Rachel voice, play/pause controls*
-  *Multilingual: 10 languages, native LLM translation*
-  *Follow-ups: Session memory, last 6 messages preserved*

How to Run

Quick Start
powershell
Start both servers
.\START.bat

Or manually:
Terminal 1: Backend
python run_server.py

Terminal 2: Frontend
cd frontend
npm run dev

Access

Frontend: <http://localhost:5173> (or 5174 if port busy)

Backend: <http://localhost:9000>

API Docs: <http://localhost:9000/docs>

📁 File Structure

Day 2 Ai Project/

```
├── backend.py      # FastAPI routes, session management
├── storyteller.py  # RAG pipeline, multimodal generation
├── document_processor.py # PDF processing, embeddings, caching
├── config.py       # Centralized configuration (models, APIs)
├── requirements.txt # Python dependencies
└── .env            # API keys (not committed)
    └── data/
        └── pdfs/      # Alice, Gulliver, Arabian Nights PDFs
    └── static/
        ├── images/    # Generated AI images
        └── audio/     # Generated audio narration
    └── frontend/
        └── src/
            ├── App.jsx   # Main component, chat logic
            ├── App.css    # ChatGPT-style UI, dark mode
            └── components/
                ├── ChatMessage.jsx # Message bubbles, media
                └── SuggestionPill.jsx # Query suggestions
            └── package.json  # Node dependencies
```

🎯 Achievements

- All core evaluation criteria (8/8)
- All bonus features (4/4)
- Professional SaaS UI (ChatGPT-style, dark mode, sidebar)
- Fast performance (sub-second retrieval, parallel generation)
- Easy model switching (centralized config)
- Open source first (90% free/OSS stack)

[System Architecture Diagram:](#)

